

Time Series Analysis - Singapore Total Fertility Rate and Total Live Births

Hunter Holt

June 08, 2026

Abstract

Over the course of the last century, Singapore's demographic transition remains as one of the most rapid ever recorded, with the Total Fertility Rate (TFR) declining from 5.76 in 1960 to less than 1.0 by 2024. This paper features a complete time series analysis of two interrelated population metrics: TFR and Total Live Births (TLB). These are scraped from the Singapore Department of Statistics' annual data, ranging from 1968 to 2024. After an established Data Analysis Lifecycle, paired with the Box-Jenkins identification cycle, a log-differencing transformation is required to achieve stationarity. This is also furthermore confirmed by ADF and KPSS tests. Autocorrelation (ACF) and partial autocorrelation (ACF) functions applied to our stationary series prompts for an $ARIMA(p, 1, q)$ candidate model(s) fitted with a logarithmic Box-Cox transform. Model selection verification using AICc identifies the specification with best fit for each series; residual diagnostics aim to replicate white noise behaviour. Out-of-sample forecasts over a subsample of 2013–2024 are evaluated against observed values, with performance assessed via RMSE, MAE, MAPE, and Theil's U. The results post-evaluation confirm the particular difficulty of accurately forecasting sub-replacement fertility prediction mapping, likely due to structural policy interventions within Singapore and/or the impact of COVID-19. This motivates additional extension processes including ARIMA-GARCH and state-space modelling.

Keywords: Total Fertility Rate; Total Live Births; ARIMA; Box-Jenkins; Stationarity; Singapore

1 Introduction

- Demographic transition
- Discuss variables
- Discuss significant change in variables, and why. EDA
- Introduce research questions: ‘which policy interventions over 1960–2024 justify the significant structural deviations in TFR and TLB?’, ‘can ARIMA models trained on the 1960–2012 window project the 2013–2024 period?’, ‘what (and how do) demographic mechanisms account for the recent significant drop in TFR?’
- State IED Box-Jenkins cycle as diagnostic tool

Singapore features one of the most clearly-defined cases of demographic transitions in the modern world. The Total Fertility Rate (TFR, defined as the average number of children a woman would bear over her lifetime given current age-specific fertility rates) dramatically decreased from 5.76 in 1960 to 0.97 in 2024, a reduction of more than 80 per cent across 64 years. Total Live Births (TLB) followed a somewhat comparable trajectory, falling from a peak of ~61,775 in 1960 to levels below 35,000 as of recent.

This pattern is partly attributable to gradual implementation of government policies. The ‘Stop-at-Two’ program, introduced in 1972, combined financial shortcomings for large families with generalised broader access to family planning (Aassve et al. 2020), resulting in the steepest decline in TFR post-1970. The subsequent government program, dubbed ‘Have Three or More if You Can Afford It’, was inevitably reversed in 1987, providing a gradual relief of TFR. More recently, the COVID-19 pandemic caused a sharp birth deferral across 2020–2022, which is an identified phenomenon within more developed economies (Aassve et al. 2020).

These features make Singapore a great investigative subject for time series modelling: the series exhibit strong trend non-stationarity, potential breaks in structure and developed patterns, and an out-of-sample period which provides more information about the series. The analysis is performed under the premise of three research questions:

1. **Policy attribution;** Which policy interventions over the documented 1960–2024 period best explain structural variation in TFR and TLB from a smooth decline trajectory?
2. **Forecast fidelity;** Can ARIMA models trained on an 80% data window (i.e. 1960–2012) accurately project the 2013–2024 period?
3. **Sub-replacement fertility;** What demographic features (age-structure shifts, ethnic differentials, and COVID-19 disruption, for instance) account for TFR falling below 1.0 recently?

The analysis featured follows the Box-Jenkins identification, estimation, and diagnostic checking cycle, as well as the Harvard Data Analysis Lifecycle (Box and Jenkins 1970).

2 Data

2.1 Data Sourcing and Structuring

The data used throughout this report are sourced from *SingaporeCensus.csv*, a population census published by the Singapore Department of Statistics. The raw comma-separated-value file is stored in wide format, i.e. the years 1960–2024 are column headers. Two rows are extracted – one for TFR and one for TLB – and reshaped to long format by using `pivot_longer()`. The resulting cleaned dataset contains 65 annual observations per series with no missing values.

For forecasting evaluation, the full series is partitioned into a training windows of $n = 53$ observations (from 1960 to 2012) and a complementary windows of $h = 12$ observations (2013 to 2024).

2.2 Statistical Summary

Table 1: Summary statistics for TFR and TLB (1960–2024).

Statistic	TFR	TLB
Min	0.97	33541.00
Q1	1.26	39570.00
Median	1.62	42484.00
Mean	2.06	44359.63
Q3	2.07	48577.00
Max	5.76	61775.00
NAs	0.00	0.00

TFR ranges from an initial post-war high of 5.76, to a final minimum of 0.97, which is much lower than the ‘replacement threshold’ of 2.1 (Espenshade, Guzman, and Westoff 2003). The considerable difference between the mean of 2.06 and median of 1.62 is reflective of a strongly right-skewed distribution produced by the pre-transitional years.

2.3 Time Series Plots




Figure 1: TFR (left) and TLB (right), from 1960–2024. The dashed vertical line indicates the train-test data partitioning.

Figure 1 indicates two salient features shared by both series. For one, a strong uniform downward trend persists over the full sample, which immediately rules out non-transformative modelling approaches. For two, the variance of the fluctuations around the trend is considerably higher during earlier years (1960–1980) instead of during the sub-replacement period post-1985. This is indicative of a multiplicative error structure rather than an additive one, which prompts for a logarithmic variance stabilisation method.

3 Context and Methodology

3.1 Transformation and Stationarity

A time series $\{X_t\}$ is said to be *weakly stationary* if its unconditional mean $\mu = E[X_t]$ is constant and its autocovariance function $\gamma(h) = \text{Cov}(X_t, X_{t-h})$ depends only on lag h , rather than time t (Box and Jenkins 1970). Both raw series are in violation of this condition; shown is a persistent downward trend (i.e. a non-constant mean) and heteroskedastic fluctuations (i.e. non-constant variance).

Since variance appears proportional to level, a logarithmic transformation $W_t = \log X_t$ is applied to stabilise variance. The log series retains a trend, prompting for a single first-difference to produce a stationary working series:

$$Y_t = (1 - B) \log X_t = \log X_t - \log X_{t-1} = \log \left(\frac{X_t}{X_{t-1}} \right),$$

where B is the backward shift operator ($BX_t = X_{t-1}$). In this context, Y_t is interpreted as the log-annual growth rate of the series. Stationarity of Y_t is assessed using two statistical tests:

- **Augmented Dickey-Fuller (ADF)** (Dickey and Fuller 1979): tests H_0 (unit root) against H_1 (stationary). A small p -value supports rejection of the unit-root null.
- **KPSS** (Kwiatkowski et al. 1992): tests H_0 : (trend-)stationary against H_1 (unit root). A large p -value (non-rejection) supports the stationarity hypothesis.

Using both tests simultaneously ensures less ambiguity that individual tests can experience; due to low power or size distortion, for instance.

3.2 ARIMA Model specification

An ARIMA(p, d, q) model applied to the log-transformed series $\log X_t$ (so that $d = 1$ acts on the log series) is:

$$\phi(B)(1 - B) \log X_t = c + \theta(B)\varepsilon_t, \quad \varepsilon_t \sim \text{WN}(0, \sigma^2),$$

where $\text{WN}(0, \sigma^2)$ denotes a white-noise sequence (an uncorrelated process with zero mean and constant variance σ^2). The AR and MA polynomials are:

$$\phi(B) = 1 - \phi_1 B - \phi_2 B^2 - \dots - \phi_p B^p, \quad \theta(B) = 1 + \theta_1 B + \theta_2 B^2 + \dots + \theta_q B^q.$$

This model is *causal* (stationary) if all roots of $\phi(B) = 0$ lie strictly outside the unit circle; additionally, it is *invertible* if all roots of $\theta(B) = 0$ lie strictly outside the unit circle.

To specify this, the `forecast` R package features `Arima(x, order = c(p, 1, q), lambda = 0)`, where `lambda = 0` applies the logarithmic Box-Cox transform internally, and automatically back-transforms forecasts to the original scale.

3.2.1 Model identification

Given a stationary series Y_t , the **autocorrelation function** (ACF) at lag k is:

$$\rho_k = \frac{\text{Cov}(Y_t, Y_{t-k})}{\text{Var}(Y_t)},$$

and the **partial autocorrelation function** (PACF) α_k measures the linear dependence between Y_t and Y_{t-k} after eliminating the effect of $Y_{t-1}, \dots, Y_{t-k+1}$. The Box-Jenkins identification table is:

Model	Theoretical ACF	Theoretical PACF
AR(p)	Geometric decay/oscillation	Cuts off after lag p
MA(q)	Decays after lag q	Geometric decay/oscillation
ARMA(p, q)	Decays after lag $q - p$	Decays after lag $p - q$

Approximate 95% bounds on the sample ACF/PACF are $\pm 1.96/\sqrt{n}$, where n is the length of the differenced series.

3.2.2 Model selection

Candidate models are compared using the corrected Akaike Information Criterion:

$$\text{AICc} = -2\hat{\ell} + 2k + \frac{2k(k+1)}{n-k-1},$$

where $\hat{\ell}$ is the maximised log-likelihood and k is the number of free parameters. AICc is typically preferred over AIC for small n (Burnham and Anderson 2002) as a stronger model complexity penalty reduces the risk of overfit and over-parameterisation.

3.3 Residual Diagnostics

An accurate, well-specified model should yield residuals $\hat{\varepsilon}_t$ that resemble white noise. This is assessed using three complementary tools:

1. **ACF of residuals.** All sample autocorrelations should lie within the 95% (i.e., $\pm 1.96/\sqrt{n}$);
2. **Ljung-Box test** (Ljung and Box 1978). The portmanteau statistic

$$Q^* = n(n+2) \sum_{k=1}^m \frac{\hat{\rho}_k^2}{n-k}$$

is compared to $\chi^2(m-p-q)$ under H_0 (i.e., no residual correlation prior to lag m). The degrees of freedom correction `fitdf` = $p + q$ accounts for parameters that are influenced by estimation error. Omitting it (setting `fitdf` = 0) artificially ‘inflates’ the χ^2 test statistic;

3. **Residual histogram:** valid Gaussian prediction intervals require residuals that are approximately Normal.

3.4 Forecast Accuracy Measurement

Point-forecast performance over our ‘hold-out’ period (2013–2024) is measured by:

$$\text{RMSE} = \sqrt{\frac{1}{h} \sum_{t=1}^h e_t^2}, \quad \text{MAE} = \frac{1}{h} \sum_{t=1}^h |e_t|, \quad \text{MAPE} = \frac{100}{h} \left| \frac{e_t}{X_t} \right|$$

where $e_t = \hat{X}_t - X_t$. As mentioned, **Theil’s U** compares model RMSE to that of the naive (that is, the random walk) forecast; $U < 1$ indicates the model outperforms the initial naive benchmark.

4 Exploratory Analysis

4.1 Log Transformation Assessment




Figure 2: Raw versus log-transformed series for TFR (top row) and TLB (bottom row); training period 1960–2012.

Figure 2 confirms that the log-transformation reduces heteroskedasticity considerably among both series; the large fluctuations earlier on in TLB are quite attenuated in comparison to the stable post-1985 period. However, a downward trend continues among both series, which prompts for first differencing.

4.2 Log-Differenced Series




Figure 3: Log-differenced series $(1 - B) \log X_t$ for TFR (left) and TLB (right), over training period 1961–2012.

4.3 Stationary Tests

4.3.1 Raw Series

Table 2: Stationarity tests on raw training series (1960–2012).

Series	ADF p-value	KPSS p-value	Conclusion
TFR	0.120	0.0100	Non-stationary
TLB	0.356	0.0205	Non-stationary

4.3.2 Log-Differenced series

Table 3: Stationarity tests on log-differenced training series.

Series	ADF p-value	Conclusion
$\Delta \log(\text{TFR})$	0.0312	Stationary
$\Delta \log(\text{TLB})$	0.1000	Stationary

4.4 ACF & PACF Analysis




Figure 4: ACF (left column) and PACF (right column) of $(1 - B) \log X_t$ for TFR (top row) and TLB (bottom row), over training period. Lags extend to 25. Blue dashed lines are 95% bounds $\pm 1.96/\sqrt{n}$.

5. Model Identification and Estimation

5.1 TFR Candidate Model

- **T1:** ARIMA(0, 1, 1); pure MA(1) on the log-differenced series;
- **T2:** ARIMA(1, 1, 1); ARMA(1, 1) on the log-differenced series;
- **T3:** ARIMA(1, 1, 0); pure AR(1) on the log-differenced series.

Table 4: Information criteria for TFR candidate models (log scale).

Model	AIC	AICc	BIC
ARIMA(0, 1, 1)	-129.171	-128.671	-123.317
ARIMA(1, 1, 1)	-127.586	-126.735	-119.781
ARIMA(1, 1, 0)	-129.137	-128.637	-123.283

The model with the lowest AICc (bolded in Table 4, for reference) is selected. Its estimated coefficients are printed below:

Table 5: Estimated parameters: ARIMA(0, 1, 1) (log-scale TFR).

Parameter	Estimate	Std. Error	z-stat	p-value
ma1	0.08201	0.15195	0.540	0.5894
drift	-0.02868	0.00988	-2.902	0.0037

5.2 TLB Candidate Model

- **L1:** ARIMA(0, 1, 1);
- **L2:** ARIMA(1, 1, 1);
- **L3:** ARIMA(0, 1, 2) (two MA terms to identify lag-2 structure).

Table 6: Information criteria for TLB candidate models (log scale).

Model	AIC	AICc	BIC
ARIMA(0, 1, 1)	-138.167	-137.667	-132.313
ARIMA(1, 1, 1)	-136.276	-135.425	-128.471
ARIMA(0, 1, 2)	-136.501	-135.650	-128.696

Table 7: Estimated parameters: ARIMA(0, 1, 1) (log-scale TLB).

Parameter	Estimate	Std.Error	z_stat	p-value
ma1	-0.10912	0.14880	-0.733	0.4634
drift	-0.00726	0.00749	-0.968	0.3328

6 Residual Diagnostics

6.1 TFR Model




Figure 5: Residual diagnostics for the selected TFR model (ARIMA(0, 1, 1)).

```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(0,1,1) with drift
## Q* = 30.331, df = 19, p-value = 0.04773
##
## Model df: 1. Total lags used: 20
```

Table 8: Ljung-Box test on TFR model residuals; lag 20, fitdf = $p + q$.

Model	Test lag	fitdf	Q*	df	p-value
ARIMA(0, 1, 1)	20	1	30.331	19	0.0477

A Ljung-Box p -value > 0.05 indicates no statistically significant residual autocorrelation in TFR up to lag 20, which supports the white-noise assumption needed for valid inference. The residual histogram (Figure 5, bottom right) should replicate a bell-shape for prediction intervals to be well-specified.

6.2 TLB Model




Figure 6: Residual diagnostics for the selected TLB model (ARIMA(0, 1, 1)).

```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,1) with drift
## Q* = 33.002, df = 19, p-value = 0.02402
##
## Model df: 1.   Total lags used: 20
```

Table 9: Ljung-Box test on TLB model residuals (lag 20, fitdf = $p + q$).

Model	Test lag	fitdf	Q*	df	p-value
ARIMA(0, 1, 1)	20	1	33.002	19	0.024

7 Forecasting

7.1 Out-of-Sample Forecasts




Figure 7: TFR forecast (2013–2024), ARIMA(0, 1, 1), with 80% and 95% prediction intervals; actual values in red.




Figure 8: TLB forecast (2013–2024), ARIMA(0, 1, 1), with 80% and 95% prediction intervals; actual values in red.

7.2 Forecasting Accuracy Evaluation

Table 10: Forecast accuracy - TFR, ARIMA(0, 1, 1) (original scale).

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	-0.0191	0.1382	0.1004	-0.2031	4.8552	0.7999	0.1577	NA
Test set	0.0424	0.0602	0.0546	3.8490	4.8760	0.4350	0.2689	1.3665

Table 11: Forecast accuracy - TLB, ARIMA(0, 1, 1) (original scale).

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1	Theil's U
Training set	-27.24307	2693.583	2030.510	-0.160790	4.528192	0.9701079	0.0051452	NA
Test set	-1712.40718	2612.449	1978.526	-4.841583	5.473900	0.9452715	0.6437107	1.923521

Report GitHub Repo

https://github.com/XCalibr3S/TSA_EDA.git

- Aassve, Arnstein, Nicolo Cavalli, Letizia Mencarini, Samuel Plach, and Massimo Livi Bacci. 2020. "The COVID-19 Pandemic and Human Fertility." *Science* 369 (6502): 370–71. <https://doi.org/10.1126/science.abc9520>.
- Box, George E. P., and Gwilym M. Jenkins. 1970. *Time Series Analysis: Forecasting and Control*. San Francisco: Holden-Day.
- Burnham, Kenneth P., and David R. Anderson. 2002. *Model Selection and Multimodel Inference: A Practical Information-Theoretic Approach*. 2nd ed. New York: Springer.
- Dickey, David A., and Wayne A. Fuller. 1979. "Distribution of the Estimators for Autoregressive Time Series with a Unit Root." *Journal of the American Statistical Association* 74 (366): 427–31. <https://doi.org/10.2307/2286348>.
- Espenshade, Thomas J., Juan Carlos Guzman, and Charles F. Westoff. 2003. "The Surprising Global Variation in Replacement Fertility." *Population Research and Policy Review* 22 (5–6): 575–83. <https://doi.org/10.1023/B:POPU.0000020882.29684.8e>.
- Kwiatkowski, Denis, Peter C. B. Phillips, Peter Schmidt, and Yongcheol Shin. 1992. "Testing the Null Hypothesis of Stationarity Against the Alternative of a Unit Root." *Journal of Econometrics* 54 (1–3): 159–78. [https://doi.org/10.1016/0304-4076\(92\)90104-Y](https://doi.org/10.1016/0304-4076(92)90104-Y).
- Ljung, Greta M., and George E. P. Box. 1978. "On a Measure of Lack of Fit in Time Series Models." *Biometrika* 65 (2): 297–303. <https://doi.org/10.1093/biomet/65.2.297>.